

Project 0: Convolutional Neural Networks

1 Objective

The ImageNet challenge initiated by Fei-Fei Li (2010) has been traditionally approached with image analysis algorithms such as SIFT with mitigated results until the late 90s. The advent of Deep Learning dawned with a breakthrough in performance which was gained by neural networks. Inspired by Yann LeCun et al. (1998) LeNet-5 model, the first deep learning model, published by Alex Krizhevsky et al. (2012) drew attention to the public by getting a top-5 error rate of 15.3% outperforming the previous best one with an accuracy of 26.2% using a SIFT model. This model, the so-called 'AlexNet', is what can be considered today as a simple architecture with five consecutive convolutional filters, max-pool layers, and three fully-connected layers.

This project is designed to provide you with first-hand experience on training a typical Convolutional Neural Network (ConvNet) model in a discriminative classification task. The model will be trained by Stochastic Gradient Descent, which is arguably the canonical optimization algorithm in Deep Learning.

2 Model

The ConvNet is a specific artificial neural network structure inspired by biological visual cortex and tailored for computer vision tasks. The ConvNet is a discriminative classifier, defined as a conditional (posterior) probability.

$$p(c \mid I; w) = \frac{\exp[f_c(I; w) + b_c]}{\sum_{c=1}^C \exp[f_c(I; w) + b_c]} \quad (1)$$

where $c \in \{1, 2, \dots, C = 10\}$ is the class label of an input image I , and $f_c(I; w)$ is the scoring function for each category, and computed by a series of operations in the ConvNet structure. Figure 1 displays an structure (architecture) of the ConvNet.

1. **Convolution.** The convolution is the core building block of a ConvNet and consists of a set of learnable filters. Every filter is a small receptive field. For example, a typical filter on the first layer of a ConvNet might have size 5x5x3 (i.e., 5 pixels width and height, and 3 color channels). Figure 2 illustrates the filter convolution.

2. **Pooling.** Pooling is a form of non-linear down-sampling. Max-pooling is the most common. It partitions the input image into a set of non-overlapping rectangles and, for each such sub-region, outputs the maximum.
3. **Relu.** Relu is non-linear activation function $f(x) = \max(0, x)$. It increases the nonlinear properties of the decision function and of the overall network without affecting the receptive fields of the convolution layer.




Figure 1: A ConvNet consists of multiple layers of filtering and sub-sampling operations for bottom-up feature extraction, resulting in multiple layers of feature maps and their sub-sampled versions. The top layer features are used for classification via multi-nomial logistic regression




Figure 2: Filter convolution. Each node denotes a 3-channel filter at specific position. Each filter spatially convolves the image.

LeNet shown in Table 2 is a typical structure design tailored for the CIFAR-10 dataset. The Block1 has 32 filters and the filter size is $5 \times 5 \times 3$. The Block5 is a logistic regression layer, in which the number of filters is the number of classes and the filter size should be the same as the output from Block4. The CIFAR-10 dataset consists of 60,000 32×32 color images in 10 classes, with 6,000 images per class. There are 50,000 training images and 10,000 testing images. Figure 3 shows example images from 10 categories. The objective of classification is to predict the category of each image.

Block1	Block2	Block3	Block4	Block5
Conv $5 \times 5 \times 3 \times 32$	Conv $5 \times 5 \times 32 \times 32$	Conv $5 \times 5 \times 32 \times 64$	Conv $4 \times 4 \times 64 \times 64$	Conv $1 \times 1 \times 64 \times 10$
Pooling 3×3	Relu	Relu	Relu	Softmax
Relu	Pooling 3×3	Pooling 3×3		

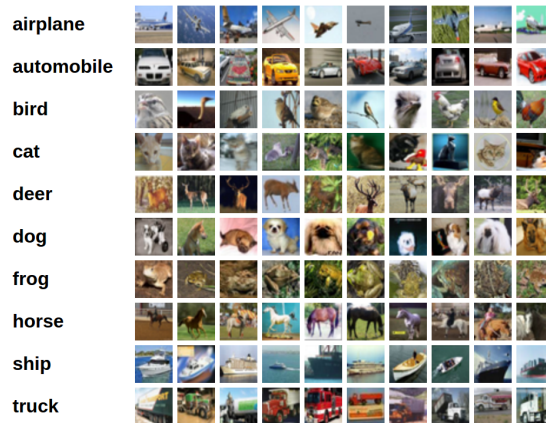


Figure 3: CIFAR-10 Dataset. CIFAR-10 contains 60,000 images in 10 classes.

3 Assignment

In this project, you will train variations of a LeNet in TensorFlow or PyTorch:

1. Optimization of LeNet.
 - (a) Complete the functions *flatten()*, *convnet_init()*, *convnet_forward()*.
 - (b) Adjust parameters *learning_rate*, *epochs* such that test-accuracy $> 70\%$.
 - (c) Plot the training loss and test accuracy over epochs in two Figures.
2. Alteration of LeNet.
 - (a) Keep the Block5, learn a ConvNet only with:
 - i. Block1.
 - ii. Block1 and Block2.
 - iii. Block1, Block2 and Block3.
 - (b) Compare the final test accuracies for (i., ii., iii.) in a Table. (**Hint:** You need to change the filter size in Block5 to match the output from previous block.)
3. Visualization of filters and activations.
 - (a) Plot the learned 32 filters of the first convolutional layer in LeNet.
 - (b) Plot the filter response maps for a given sample image of CIFAR-10.

Submission: (1) Distill your results into a report as pdf, and, (2) zip your code.